









Test setup: All VM's pinned to the same CPU. Equal load on all vCPUs







Test setup: 50% of VM's pinned to pCPU 0 and the rest to pCPU 1. Equal load on all vCPUs







Test setup: Balanced pinning of vCPU to pCPU. Equal load on all vCPUs









Test setup: Each vCPU has equal affinity to all pCPUs. Equal load on all vCPUs







Test setup: Each vCPU has equal affinity to all pCPUs. 50% vCPUs run heavy load and the remaining run a light load

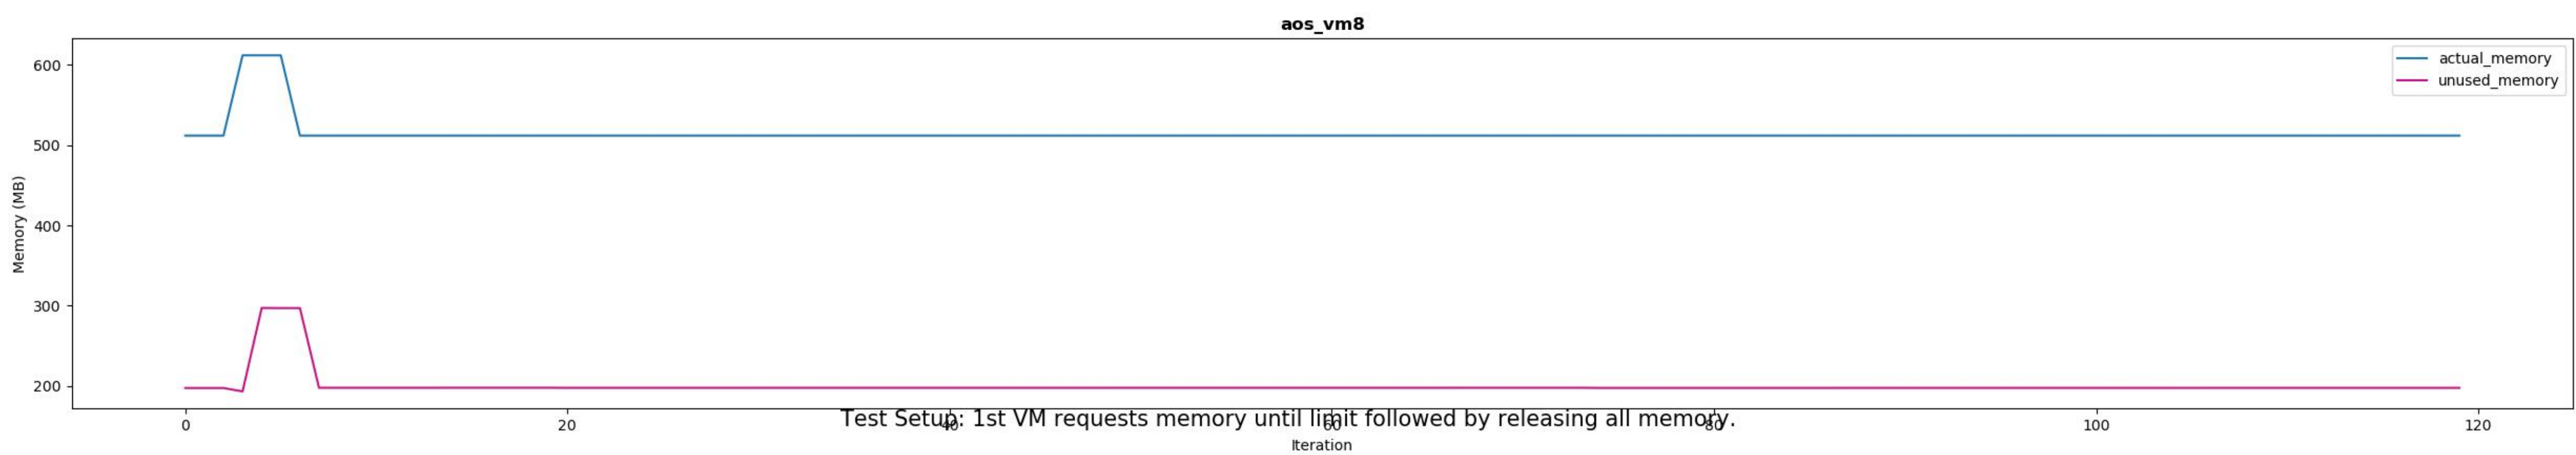





Test setup: Balanced pinning of vCPU to pCPU. No load on any vCPUs





Test setup: Each vCPU has equal affinity to all pCPUs. All vCPUs run light load, 50% switch to heavy load later



Test Setup: 1st VM requests memory until limit followed by releasing all memory.



